

python_data_science_intro

March 20, 2019

```
In [1]: ### AUTHOR: Tim Meehan  
### PRESENTED: python hack night, 2/19/19  
### contact: on slack at hack-portland.slack.com
```

```
In [2]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
%matplotlib inline
```

1 Pandas

pandas is a python library that provides a dataframe structure (think excel spreadsheet). It contains built in data manipulation and statistics functions. This notebook will use dataframes built from python dicts and read from csv's. It has a built in read function for many formats, see <http://pandas.pydata.org/pandas-docs/version/0.15/io.html> for complete list.

```
In [3]: # Example 1 - creating a dataframe
```

```
In [4]: # create a new dataframe and store it under 'dataframe'  
dataframe = pd.DataFrame(  
    {'col1': [1, 2, 3, 4, 5, 6],  
     'col2': ['a', 'b', 'b', 'a', 'a', 'b'],  
     'col3': [130, 145, 145, 123, 140, 157]})
```

```
In [5]: # print the dataframe to the output  
dataframe.head()
```

```
Out[5]:  col1 col2 col3  
0  1  a  130  
1  2  b  145  
2  3  b  145  
3  4  a  123  
4  5  a  140
```

```
In [6]: # count the number of times a value occurs in a column  
dataframe['col2'].value_counts()
```

```
Out[6]: b 3
        a 3
        Name: col2, dtype: int64
```

```
In [7]: # filter the dataframe based on column values
        dataframe[dataframe['col3'].gt(140) & dataframe['col1'].gt(2)]
```

```
Out[7]:  col1 col2 col3
        2 3 b 145
        5 6 b 157
```

```
In [8]: # a filter just takes a boolean mask as input
        dataframe['col3'].gt(140) & dataframe['col1'].gt(2)
```

```
Out[8]: 0 False
        1 False
        2 True
        3 False
        4 False
        5 True
        dtype: bool
```

```
In [9]: # add a column, that is the sum of two columns
        dataframe['col4'] = dataframe['col1'] + dataframe['col3']
        dataframe
```

```
Out[9]:  col1 col2 col3 col4
        0 1 a 130 131
        1 2 b 145 147
        2 3 b 145 148
        3 4 a 123 127
        4 5 a 140 145
        5 6 b 157 163
```

```
In [10]: # define some complicated function, and apply it to each value in a column (series)
        def mod_col_3(val):
            """if val gt 140 return 1, else return 0"""
            if val > 140:
                output = 1
            else:
                output = 0
            return output

        # the apply function returns a new series, it will not change the input series
        # we will assign the result to a new column in the df
        dataframe['col5'] = dataframe['col3'].apply(mod_col_3)
        dataframe
```

```
Out[10]:  col1 col2 col3 col4 col5
        0 1 a 130 131 0
```

```

1 2 b 145 147 1
2 3 b 145 148 1
3 4 a 123 127 0
4 5 a 140 145 0
5 6 b 157 163 1

```

```

In [11]: # get statistics automatically on all numeric columns
dataframe.describe()

```

```

Out[11]:  col1  col3  col4  col5
count  6.000000  6.000000  6.000000  6.000000
mean    3.500000  140.000000  143.500000  0.500000
std     1.870829  12.066483  12.988456  0.547723
min     1.000000  123.000000  127.000000  0.000000
25%     2.250000  132.500000  134.500000  0.000000
50%     3.500000  142.500000  146.000000  0.500000
75%     4.750000  145.000000  147.750000  1.000000
max     6.000000  157.000000  163.000000  1.000000

```

```

In [12]: #take the mean of col3
dataframe.col3.mean()

```

```

Out[12]: 140.0

```

2 matplotlib

matplotlib is the most standard plotting library for python, however there are many more to chose from.

```

In [13]: # pandas has some matplotlib functionality built in

```

```

In [14]: dataframe

```

```

Out[14]:  col1  col2  col3  col4  col5
0 1 a 130 131 0
1 2 b 145 147 1
2 3 b 145 148 1
3 4 a 123 127 0
4 5 a 140 145 0
5 6 b 157 163 1

```

```


In [15]: # example bar plot
dataframe.plot(x='col1', y='col3', kind='bar', color='b')

```

```


Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x118f2a390>

```




```
In [16]: dataframe['col2'].value_counts().plot(kind='bar', color='b')
```

```
Out[16]: <matplotlib.axes._subplots.AxesSubplot at 0x11b625550>
```




```
In [17]: # example scatter plot
         dataframe.plot('col1', 'col3', kind='scatter', rot=90)
```

```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x11b6a7e10>
```



```
In [18]: # dual y-axis
         ax = dataframe.plot('col1', 'col3', kind='scatter', label='col3')
         ax2 = ax.twinx() #pyplot
         dataframe.plot('col1', 'col4', kind='scatter', ax=ax2, label='col4', color='orange')
         ax2.legend(loc=1) #pyplot
```

```
Out[18]: <matplotlib.legend.Legend at 0x11b848a58>
```



2.1 data from csv

In [19]: *### this section will introduce reading csv, groupby, merge and pivot_table functions*

In [20]: `iris = pd.read_csv('https://raw.githubusercontent.com/mwaskom/seaborn-data/master/iris.csv')`

In [21]: `iris.shape`


Out[21]: (150, 5)

In [22]: `iris.head()`

```
Out[22]:  sepal_length  sepal_width  petal_length  petal_width  species
0      5.1      3.5      1.4      0.2  setosa
1      4.9      3.0      1.4      0.2  setosa
2      4.7      3.2      1.3      0.2  setosa
3      4.6      3.1      1.5      0.2  setosa
4      5.0      3.6      1.4      0.2  setosa
```

In [23]: *# check avg of each value based on species*
`iris.groupby('species').mean().plot(kind='bar')`

Out[23]: <matplotlib.axes._subplots.AxesSubplot at 0x11b89ddd8>



```
In [24]: grouped = iris.groupby('species').agg({'sepal_length': ['mean', 'std'],
        'sepal_width': 'count'})
```

```
In [25]: grouped
```


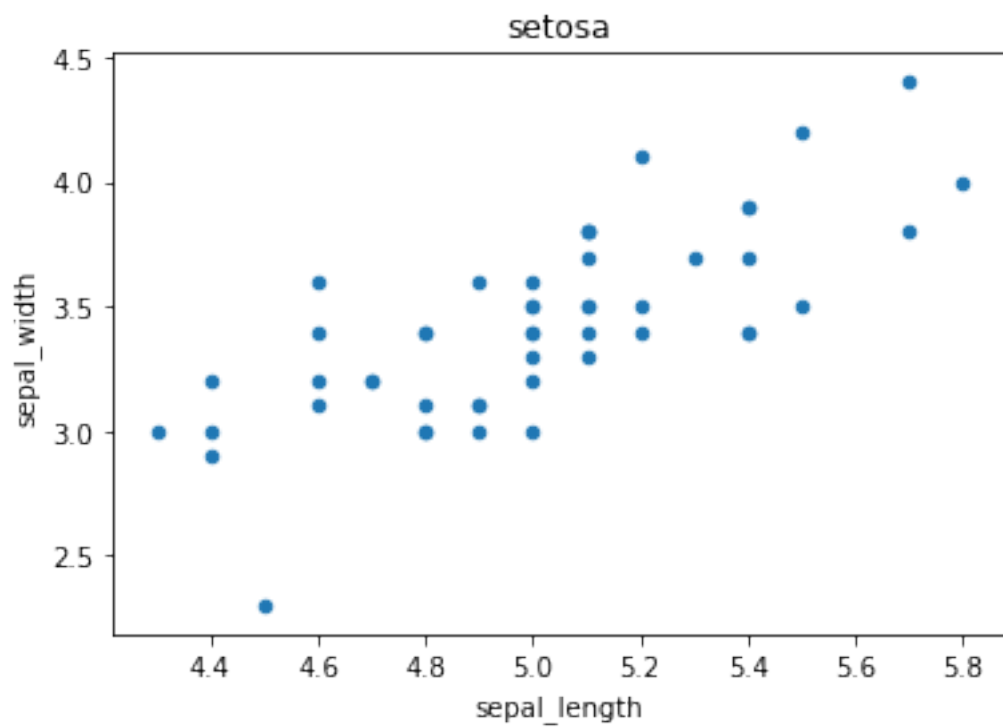
```
Out[25]: sepal_length sepal_width
        mean std count
species
setosa  5.006  0.352490  50
versicolor  5.936  0.516171  50
virginica  6.588  0.635880  50
```


```
In [26]: grouped.columns.values
```

```
Out[26]: array([(('sepal_length', 'mean'), ('sepal_length', 'std')),
        ('sepal_width', 'count')], dtype=object)
```

```
In [27]: grouped.columns = ['_'.join(col).strip() for col in grouped.columns.values]
```

```
In [28]: # iterate through subset of dataframe
        # each "df" in forloop is a filtered dataframe based on the groupby argument
        # "grp" will contain the column val, or list of values, that the current df is filter
        for grp, df in iris.groupby('species'):
            df.plot('sepal_length', 'sepal_width', kind='scatter', title=grp)
```





```
In [29]: exercise = pd.read_csv('https://raw.githubusercontent.com/mwaskom/seaborn-data/master/
```

```
In [30]: exercise.head()
```

```
Out[30]:  id diet pulse time kind
0  1 low fat 85 1 min rest
1  1 low fat 85 15 min rest
2  1 low fat 88 30 min rest
3  2 low fat 90 1 min rest
4  2 low fat 92 15 min rest
```

```
In [31]: exercise['time'].unique()
```

```
Out[31]: array(['1 min', '15 min', '30 min'], dtype=object)
```

```
In [32]: exercise.pivot_table(index=['diet','kind'], columns='time', values='pulse', aggfunc='r
```

```
Out[32]: time 1 min 15 min 30 min
diet kind
low fat rest 97 97 94
         running 98 132 120
         walking 95 103 104
no fat rest 100 99 100
         running 103 135 150
         walking 103 109 103
```

```
In [33]: df1 = pd.DataFrame({'employee': ['Bob', 'Jake', 'Lisa', 'Sue'],
    'group': ['Accounting', 'Engineering', 'Engineering', 'HR']})
df2 = pd.DataFrame({'employee': ['Lisa', 'Bob', 'Jake', 'Sue'],
    'hire_date': [2004, 2008, 2012, 2014]})
```

```
In [34]: df1
```

```
Out[34]:  employee group
0  Bob  Accounting
1  Jake  Engineering
2  Lisa  Engineering
3  Sue   HR
```

```
In [35]: df2
```

```
Out[35]:  employee hire_date
0  Lisa  2004
1  Bob   2008
2  Jake  2012
3  Sue   2014
```

```
In [36]: df1.merge(df2, on=['employee'], how='left')
```

```
Out[36]:  employee group hire_date
0  Bob  Accounting  2008
1  Jake  Engineering 2012
2  Lisa  Engineering 2004
3  Sue   HR         2014
```

2.2 Putting it all together

```
In [37]: ### In this example, we look at (fake) test results
### we will create fake date labels, find the number of passes and fails in each,
### then average the fails by day, which will give you yield (pass/fail is 1,0)
### In addition, for those widgets that passed or failed, we will plot a box plot
### to show the distribution of some parameter, could be temperature, time
### completing some process step etc...

### the end chart will show by day, how many parts were tested, the yield of those pa
### and the distribution of some parameter, so you can see the change over time and r
### to the parameter
```

```
In [38]: ### create a fake time index. Use pd.Series to create a single column
date_index = pd.Series(['2018-03-03']*10 + ['2018-03-04']*12 + ['2018-03-05']*18 + ['2018-03-06']*10)
date_index = pd.to_datetime(date_index)
# print out preview of what the "column" looks like
date_index.head()
```

```

Out [38]: 0 2018-03-03
          1 2018-03-03
          2 2018-03-03
          3 2018-03-03
          4 2018-03-03
          Name: date, dtype: datetime64[ns]

In [39]: ### use random numbers to generate pass and fail data
         from random import choices, gauss
         test1_yld = pd.Series(choices(population=[0, 1], weights=[30, 70], k=91), name='test1_pass_fail')
         test2_yld = pd.Series(choices([0, 1], [10, 90], k=91), name='test2_pass_fail')

         ### create fake process parameter. Think seconds a widget spent at process step x
         process_param = pd.Series([gauss(20, 4) for i in range(91)], name='process_param')

In [40]: ### we will also create a fake "widget_id"
         widget_id = pd.Series([i for i in range(len(date_index))], name='widget_id')

In [41]: ### concat the fail data together, think of the python zip() function
         ### the resulting dataframe will have for each date, a number of passes and fails
         ### think of this as a single widget per row
         example_df = pd.concat([date_index, test1_yld, test2_yld, process_param, widget_id], axis=1)
         example_df.head()

Out [41]:  date test1_pass_fail test2_pass_fail process_param widget_id
          0 2018-03-03 1 1 16.154771 0
          1 2018-03-03 1 1 23.180473 1
          2 2018-03-03 1 1 17.247763 2
          3 2018-03-03 1 1 19.270181 3
          4 2018-03-03 1 1 28.049684 4

In [42]: ### summarize the average pass value by date. This will be your yield trend
         ### you can also aggregate the "nunique" pandas function to count the unique number of widgets
         yield_trend = example_df.groupby('date').agg({'test1_pass_fail': 'mean',
              'test2_pass_fail': 'mean',
              'widget_id': lambda rows: rows.nunique()}) #rows is the *grouped* series of data

         # rename the columns
         yield_trend.rename(columns = {'test1_pass_fail': 'test1_yld',
              'test2_pass_fail': 'test2_yld',
              'widget_id': 'num_widgets'},
              inplace=True)

         # print the dataframe
         yield_trend

Out [42]:  test1_yld test2_yld num_widgets
          date
2018-03-03  0.600000  1.000000      10

```

```

2018-03-04 0.666667 0.833333 12
2018-03-05 0.833333 1.000000 18
2018-03-06 0.571429 0.714286 7
2018-03-07 0.888889 0.777778 9
2018-03-08 0.500000 0.900000 20
2018-03-09 0.800000 0.933333 15

```

```

In [43]: # set up global plot size
from pylab import rcParams
rcParams['figure.figsize'] = 13, 7

In [44]: # plot bar chart for widget count in graph (1st y axis)
ax = yield_trend.plot(y='num_widgets', kind='bar', alpha=0.4, color='r', legend=None)

# for each bar, put a text label above it with the actual count value
for p in ax.patches:
    ax.annotate(np.round(p.get_height(), decimals=2), (p.get_x()+p.get_width()/2., p.get_y()+p.get_height()/2.))

# establish your 2nd and 3rd y axis and positions
ax2 = ax.twinx()
ax3 = ax.twinx()
ax3.spines["right"].set_position(("axes", 1.1))

# on your 3rd y axis, plot a boxplot. This will show for each date, this distribution
bp = example_df.boxplot(column='process_param', by='date', ax=ax3, widths = 0.25, return_type='dict')

# for each yield column, make a line plot, your label will be you column in the legend
for col in ['test1_yld', 'test2_yld']:
    ax2.plot(ax.get_xticks(), yield_trend[col], label=col, marker='o');
ax2.legend(loc='lower center', bbox_to_anchor=(0.5, 0.90),
           fancybox=True, shadow=True, ncol=5);

plt.suptitle('Yield Trend Chart')
plt.title('')

ax2.grid(True, axis='y')
ax.grid(True, axis='x')
ax3.grid(False)
ax.set_ylabel('WIDGET COUNT')
ax2.set_ylabel('YIELD')
ax3.set_ylabel('seconds at process step x')
ax.set_ylim(0, 40);

len(plt.gca().get_xticks());
# plt.show();

```

